S (3.2M): learning excess at α = 1

M (25M): learning excess at α = 1

L (85M): learning excess at α = 1

ε_learn, bits/token (dashed: ε* = 0.02)

training tokens T

n = 8
n = 16
n = 32
n = 64
n = 128

complexity multiplier M_K shrinks with size (G6: L/S = 0.39)

data multiplier at ε* (log)

model-free (G1, n = 128 vs 8, α = 1)
law, M_K_both fit rows (G6)

S (3.2M)
M (25M)
L (85M)

both elasticities fall with size; alignment's falls faster

elasticity of T* (at ε*, K_ref)

e_K (complexity)
e_α (alignment)

E = 0.80
E = 0.47
E = 0.36

alignment share r_α of the excess (G5)

early positions (1-16)
late positions (65-128)